

In [62]:

```
import numpy as np
import pandas as pd
```

In [63]:

```
path_to_dataset = '/content/drive/MyDrive/Colab Notebooks/project-ea/bitstampUSD_1-min_data_2012-01-01_to_2021-03-31.csv'
bitcoin_dataset = pd.read_csv(path_to_dataset)
```

In [64]:

```
#showing the data
bitcoin_dataset
```

Out[64]:

| | Timestamp | Open | High | Low | Close | Volume_(BTC) | Volume_(Currency) | Weighted_Price |
|---------|------------|----------|----------|----------|----------|--------------|-------------------|----------------|
| 0 | 1325317920 | 4.39 | 4.39 | 4.39 | 4.39 | 0.455581 | 2.000000 | 4.390000 |
| 1 | 1325317980 | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| 2 | 1325318040 | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| 3 | 1325318100 | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| 4 | 1325318160 | NaN | NaN | NaN | NaN | NaN | NaN | NaN |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 4857372 | 1617148560 | 58714.31 | 58714.31 | 58686.00 | 58686.00 | 1.384487 | 81259.372187 | 58692.753339 |
| 4857373 | 1617148620 | 58683.97 | 58693.43 | 58683.97 | 58685.81 | 7.294848 | 428158.146640 | 58693.226508 |
| 4857374 | 1617148680 | 58693.43 | 58723.84 | 58693.43 | 58723.84 | 1.705682 | 100117.070370 | 58696.198496 |
| 4857375 | 1617148740 | 58742.18 | 58770.38 | 58742.18 | 58760.59 | 0.720415 | 42332.958633 | 58761.866202 |
| 4857376 | 1617148800 | 58767.75 | 58778.18 | 58755.97 | 58778.18 | 2.712831 | 159417.751000 | 58764.349363 |

4857377 rows × 8 columns

In [65]:

```
bitcoin_dataset.describe()
```

Out[65]:

| | Timestamp | Open | High | Low | Close | Volume_(BTC) | Volume_(Currency) | Weighted_P |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|-------------------|--------------|
| count | 4.857377e+06 | 3.613769e+06 | 3.613769e+06 | 3.613769e+06 | 3.613769e+06 | 3.613769e+06 | 3.613769e+06 | 3.613769e+06 |
| mean | 1.471301e+09 | 6.009024e+03 | 6.013357e+03 | 6.004488e+03 | 6.009014e+03 | 9.323249e+00 | 4.176284e+04 | 6.008935e+03 |
| std | 8.428019e+07 | 8.996247e+03 | 9.003521e+03 | 8.988778e+03 | 8.996360e+03 | 3.054989e+01 | 1.518248e+05 | 8.995992e+03 |
| min | 1.325318e+09 | 3.800000e+00 | 3.800000e+00 | 1.500000e+00 | 1.500000e+00 | 0.000000e+00 | 0.000000e+00 | 3.800000e+00 |
| 25% | 1.398179e+09 | 4.438600e+02 | 4.440000e+02 | 4.435200e+02 | 4.438600e+02 | 4.097759e-01 | 4.521422e+02 | 4.438306e+02 |
| 50% | 1.471428e+09 | 3.596970e+03 | 3.598190e+03 | 3.595620e+03 | 3.597000e+03 | 1.979811e+00 | 3.810124e+03 | 3.596804e+03 |
| 75% | 1.544288e+09 | 8.627270e+03 | 8.632980e+03 | 8.621090e+03 | 8.627160e+03 | 7.278216e+00 | 2.569821e+04 | 8.627637e+03 |
| max | 1.617149e+09 | 6.176356e+04 | 6.178183e+04 | 6.167355e+04 | 6.178180e+04 | 5.853852e+03 | 1.390067e+07 | 6.171621e+04 |

In [6]:

```
#checking for nulls
bitcoin_dataset.isnull().sum()
```

Out[6]:

```
Out[7]:
Timestamp          0
Open              1243608
High              1243608
Low               1243608
Close             1243608
Volume_(BTC)      1243608
Volume_(Currency) 1243608
Weighted_Price    1243608
dtype: int64
```

In [8]:

```
#we need to know about how much percent of data is required for proper dataset
bitcoin_dataset.isnull().mean().round(4) * 100
```

Out[8]:

```
Timestamp          0.0
Open              25.6
High              25.6
Low               25.6
Close             25.6
Volume_(BTC)      25.6
Volume_(Currency) 25.6
Weighted_Price    25.6
dtype: float64
```

In [9]:

```
#in each row we have like 28% null values. Even though it's lot of null data, it's still
invalid and so safe to delete!!!
bitcoin_dataset.dropna(inplace=True)
```

In [10]:

```
#now we can inspect the null free dataser!!
bitcoin_dataset
```

Out[10]:

| | Timestamp | Open | High | Low | Close | Volume_(BTC) | Volume_(Currency) | Weighted_Price |
|---------|------------|----------|----------|----------|----------|--------------|-------------------|----------------|
| 0 | 1325317920 | 4.39 | 4.39 | 4.39 | 4.39 | 0.455581 | 2.000000 | 4.390000 |
| 478 | 1325346600 | 4.39 | 4.39 | 4.39 | 4.39 | 48.000000 | 210.720000 | 4.390000 |
| 547 | 1325350740 | 4.50 | 4.57 | 4.50 | 4.57 | 37.862297 | 171.380338 | 4.526411 |
| 548 | 1325350800 | 4.58 | 4.58 | 4.58 | 4.58 | 9.000000 | 41.220000 | 4.580000 |
| 1224 | 1325391360 | 4.58 | 4.58 | 4.58 | 4.58 | 1.502000 | 6.879160 | 4.580000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 4857372 | 1617148560 | 58714.31 | 58714.31 | 58686.00 | 58686.00 | 1.384487 | 81259.372187 | 58692.753339 |
| 4857373 | 1617148620 | 58683.97 | 58693.43 | 58683.97 | 58685.81 | 7.294848 | 428158.146640 | 58693.226508 |
| 4857374 | 1617148680 | 58693.43 | 58723.84 | 58693.43 | 58723.84 | 1.705682 | 100117.070370 | 58696.198496 |
| 4857375 | 1617148740 | 58742.18 | 58770.38 | 58742.18 | 58760.59 | 0.720415 | 42332.958633 | 58761.866202 |
| 4857376 | 1617148800 | 58767.75 | 58778.18 | 58755.97 | 58778.18 | 2.712831 | 159417.751000 | 58764.349363 |

3613769 rows x 8 columns

In [11]:


```
#for more visualization, we need to import seaborn and matplotlib
import seaborn as sbn
import matplotlib.pyplot as plt
```

In [12]:

```
#correlation helps us to find out which of the fields are related to each other..
plt.figure(figsize=(15, 15))
sbn.heatmap(bitcoin_dataset.corr(), annot=True, cmap='rainbow', linewidths=1, linecolor='black')
```

Out[12]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fcb89d35d50>



In [13]:

```
#we should change the names of column just in case to avoid any feature errors
bitcoin_dataset.rename(columns={"Volume_(BTC)" : "Volume_BTC", "Volume_(Currency)" : "Volume_Currency"}, inplace=True)
```

In [14]:

```
#similarly, we need to get the timestamp values to be human understandable
```

```
bitcoin_dataset['New_Dates'] = pd.to_datetime(bitcoin_dataset['Timestamp'], unit='s')

bitcoin_dataset
```

Out[14]:

| | Timestamp | Open | High | Low | Close | Volume_BTC | Volume_Currency | Weighted_Price | New_Dates |
|---------|------------|----------|----------|----------|----------|------------|-----------------|----------------|---------------------|
| 0 | 1325317920 | 4.39 | 4.39 | 4.39 | 4.39 | 0.455581 | 2.000000 | 4.390000 | 2011-12-31 07:52:00 |
| 478 | 1325346600 | 4.39 | 4.39 | 4.39 | 4.39 | 48.000000 | 210.720000 | 4.390000 | 2011-12-31 15:50:00 |
| 547 | 1325350740 | 4.50 | 4.57 | 4.50 | 4.57 | 37.862297 | 171.380338 | 4.526411 | 2011-12-31 16:59:00 |
| 548 | 1325350800 | 4.58 | 4.58 | 4.58 | 4.58 | 9.000000 | 41.220000 | 4.580000 | 2011-12-31 17:00:00 |
| 1224 | 1325391360 | 4.58 | 4.58 | 4.58 | 4.58 | 1.502000 | 6.879160 | 4.580000 | 2012-01-01 04:16:00 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 4857372 | 1617148560 | 58714.31 | 58714.31 | 58686.00 | 58686.00 | 1.384487 | 81259.372187 | 58692.753339 | 2021-03-30 23:56:00 |
| 4857373 | 1617148620 | 58683.97 | 58693.43 | 58683.97 | 58685.81 | 7.294848 | 428158.146640 | 58693.226508 | 2021-03-30 23:57:00 |
| 4857374 | 1617148680 | 58693.43 | 58723.84 | 58693.43 | 58723.84 | 1.705682 | 100117.070370 | 58696.198496 | 2021-03-30 23:58:00 |
| 4857375 | 1617148740 | 58742.18 | 58770.38 | 58742.18 | 58760.59 | 0.720415 | 42332.958633 | 58761.866202 | 2021-03-30 23:59:00 |
| 4857376 | 1617148800 | 58767.75 | 58778.18 | 58755.97 | 58778.18 | 2.712831 | 159417.751000 | 58764.349363 | 2021-03-31 00:00:00 |

3613769 rows x 9 columns

In [15]:


```
required_features = ['Open', 'High', 'Low', 'Volume_BTC', 'Volume_Currency', 'Weighted_P
rice']
output_label = 'Close'
```


In []:

```
bitcoin_dataset
```

In [17]:

```
bitcoin_dataset.hist(bins=20, legend=False, figsize=(25, 20))
plt.show()
```






In [20]:

```
import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(12,8))
sns.scatterplot(x='Low', y='High', data=bitcoin_dataset)
plt.xlabel("x-axis (Low Price)")
plt.ylabel("y-axis (High Price)")
plt.legend()
plt.show()
```


WARNING:matplotlib.legend:No handles with labels found to put in legend.



In [22]:

```
plt.figure(figsize=(12,8))
sns.scatterplot(x='Open', y='Close', data=bitcoin_dataset)
plt.xlabel("x-axis (Open Price)")
plt.ylabel("y-axis (Close Price)")
plt.legend()
plt.show()
```

WARNING:matplotlib.legend:No handles with labels found to put in legend.




In [24]:

```
bitcoin_dataset.dtypes.value_counts().plot.pie(autopct='%0.2f%%')
```

Out[24]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fcb86023210>



In [27]:

```
all_cols = bitcoin_dataset.select_dtypes(include=('float','int')).columns
float_data = bitcoin_dataset.select_dtypes(include=('float')).columns
float_data
```

Out[27]:


```
Index(['Open', 'High', 'Low', 'Close', 'Volume_BTC', 'Volume_Currency',
      'Weighted_Price'],
      dtype='object')
```

In [34]:

```
plt.figure(figsize=(12,9))
sns.heatmap(bitcoin_dataset.isnull())
```

Out[34]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fcb863b4c10>



In [37]:

```
def plotfig(data):
    plt.figure(figsize = (16,8))
    plt.plot(data["Weighted_Price"])
    plt.xlabel('Time',fontsize=15)
    plt.ylabel('Weighted_Price',fontsize=15)
    plt.xticks(rotation=45)
    plt.show()
```

In [39]:

```
bitcoin_dataset["Timestamp"] = pd.to_datetime(bitcoin_dataset["Timestamp"], infer_datetime_format=True, unit="s")
data = bitcoin_dataset.set_index("Timestamp")

# Considering data of last 70 days only
freq = 60
c = int(60/freq)
days = 70


data = data.tail(days*24*60)
plotfig(data)
print(data.shape)

#taking interval of 10 minutes

data = data[::freq]
data.dropna(axis=0,inplace=True)

data = data["Weighted_Price"]

data = data.values.reshape(-1,1)
print("No. of days: ",data.shape[0]/(24*c))
```



(100800, 8)
No. of days: 70.0

```
/usr/local/lib/python3.7/dist-packages/pandas/util/_decorators.py:311: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
return func(*args, **kwargs)

In [41]:

```
#scaling
from sklearn.preprocessing import MinMaxScaler
scaler_train = MinMaxScaler(feature_range=(0, 1))
data = scaler_train.fit_transform(data)
```

In [42]:

```
# Train-test split

train_split_time = 30*24*c
test_split_time = 50*24*c

train_data = data[:train_split_time]
valid_data = data[train_split_time:test_split_time]
test_data = data[test_split_time:]

window_size = 30
batch_size = 32
shuffle_buffer_size = 1000
```

In [43]:

```
data.shape
```

Out[43]:

```
(1680, 1)
```

In [50]:

```
Total = ['Open', 'High', 'Low', 'Close', 'Volume_BTC', 'Volume_Currency', 'Weighted_Price']
btc_new_x = ['Open', 'High', 'Low', 'Volume_BTC', 'Volume_Currency', 'Weighted_Price']
```



```
btc_new_y = ['Close']
```

In [52]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split( bitcoin_dataset[btc_new_x], bitcoin_dataset[btc_new_y], test_size = 0.3)
```

In [47]:

```
from sklearn.linear_model import LinearRegression
```

In [53]:

```
regression_model = LinearRegression()
regression_model.fit(X_train, y_train)
```

Out[53]:

```
LinearRegression()
```

In [54]:

```
#our model is ready!! Time to test accuracy!!!
regression_model.score(X_test, y_test)
```

Out[54]:

```
1.0
```

In [56]:

```
# We got 99% accuracy on our test data also..that means our model is quite good
#now we can test for actual predictions!!

#we will take some data from test set and try to predict that

#we will take row number 55 from our original dataset, i.e from one before splitting
sample_data = bitcoin_dataset.iloc[55]
sample_data
```

Out[56]:

```
Timestamp          2012-01-06 18:42:00
Open                6.4
High                6.4
Low                 6.4
Close               6.4
Volume_BTC          9.110853
Volume_Currency     58.309457
Weighted_Price       6.4
New_Dates            2012-01-06 18:42:00
Name: 9290, dtype: object
```

from above, it's clear that original price was 6.4 and out model predicted it as 6.4 which is pretty much equal

In [58]:

```
#to make predictions of future values we will need to shift data by 30 days!!
future_set = bitcoin_dataset.shift(periods=30).tail(30)
```

In [59]:

```
#similarly we can use r2_score to see our accuracy
from sklearn.metrics import r2_score
```

In [60]:

```
predictions = regression_model.predict(X_test)
print('Accuracy of model : ', r2_score(predictions, y_test))
```

```
Accuracy of model : 1.0
```

